

```
In [1]: from IPython.display import Image
Image(filename = "C:/Users/HP/Downloads/auto.png")
```

Out[1]:



Auto MPG

Donated on 7/6/1993

Revised from CMU StatLib library, data concerns city-cycle fuel consumption

Dataset Characteristics

Multivariate

Subject Area

Other

Associated Tasks

Regression

Feature Type

Real, Categorical, Integer

Instances

398

Features

7

Dataset Information

Additional Information

This dataset is a slightly modified version of the dataset provided in the StatLib library. In line with the use by Ross Quinlan (1993) in predicting the attribute "mpg", 8 of the original instances were removed because they had unknown values for the "mpg" attribute. The original dataset is available in the file "auto-mpg.data-original"....

SHOW MORE ▾

Import Basic Packages

```
In [2]: import os
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set()

import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: file_path = "C:/Users/HP/Downloads/mpg.txt"
with open(file_path, "r") as file:
    datadescription = file.read()
print(datadescription)
```

1. Title: Auto-Mpg Data
2. Sources:
 - (a) Origin: This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University. The dataset was used in the 1983 American Statistical Association Exposition.
 - (c) Date: July 7, 1993
3. Past Usage:
 - See 2b (above)
 - Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan

4. Relevant Information:

This dataset is a slightly modified version of the dataset provided in the StatLib library. In line with the use by Ross Quinlan (1993) in predicting the attribute "mpg", 8 of the original instances were removed because they had unknown values for the "mpg" attribute. The original dataset is available in the file "auto-mpg.data-original".

"The data concerns city-cycle fuel consumption in miles per gallon, to be predicted in terms of 3 multivalued discrete and 5 continuous attributes." (Quinlan, 1993)

5. Number of Instances: 398

6. Number of Attributes: 9 including the class attribute

7. Attribute Information:

1. mpg: continuous
2. cylinders: multi-valued discrete
3. displacement: continuous
4. horsepower: continuous
5. weight: continuous
6. acceleration: continuous
7. model year: multi-valued discrete
8. origin: multi-valued discrete
9. car name: string (unique for each instance)

8. Missing Attribute Values: horsepower has 6 missing values

```
In [4]: data = pd.read_csv('mpg.csv')
data.head()
```

```
Out[4]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin	name
0	18.0	8	307.0	130	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	ford torino

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              398 non-null   float64
1   cylinders        398 non-null   int64
2   displacement     398 non-null   float64
3   horsepower       398 non-null   object
4   weight           398 non-null   int64
5   acceleration     398 non-null   float64
6   model_year      398 non-null   int64
7   origin           398 non-null   int64
8   name             398 non-null   object
```

```
dtypes: float64(3), int64(4), object(2)
memory usage: 28.1+ KB
```

No missing values


Handling char val

As horsepower is numerical value but here in dataset it is given object .. so we will convert this column into int

```
In [6]: def horsepwr(h):
        h = h.replace('?', '0')
        h = int(h)
        return h
data['horsepower_new'] = data['horsepower'].map(horsepwr)
data.drop('horsepower', axis=1, inplace = True)
```

```
In [7]: sns.boxplot(data['horsepower_new'])
```

```
Out[7]: <Axes: >
```



```
In [8]: data['horsepower_new'] = data['horsepower_new'].replace(0, data['horsepower_new'].median)
```


```
In [9]: data.isnull().sum()
```



```
Out[9]: mpg           0
        cylinders     0
        displacement  0
        weight        0
```

```
acceleration      0
model_year        0
origin            0
name              0
horsepower_new    0
dtype: int64
```


Treating outlier

```
In [10]: def boxplot(col):
          sns.boxplot(data[col])
          plt.show()
          for i in list(data.select_dtypes(exclude = ['object']).columns)[0:]:
              boxplot(i)
```






5000
4500
4000
3500
3000
2500
2000
1500






0

25.0
22.5
20.0
17.5
15.0
12.5
10.0
7.5



0






In [11]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   mpg                   398 non-null   float64
1   cylinders              398 non-null   int64
2   displacement           398 non-null   float64
3   weight                 398 non-null   int64
4   acceleration           398 non-null   float64
5   model_year            398 non-null   int64
6   origin                 398 non-null   int64
7   name                   398 non-null   object
8   horsepower_new         398 non-null   int64
dtypes: float64(3), int64(5), object(1)
memory usage: 28.1+ KB
```

In [12]: `sns.boxplot(data['acceleration'])`

Out[12]: `<Axes: >`



Outlier Treatment

```
In [13]: Q1 = data['acceleration'].quantile(0.25)
Q3 = data['acceleration'].quantile(0.75)
IQR = Q3 - Q1
upper_limit = Q3 + 1.5*IQR
lower_limit = Q1 - 1.5*IQR
print('Q1: ',Q1)
print("Q3: ",Q3)
print('IQR: ',IQR)
print('upper limit: ',upper_limit)
print('lower limit: ',lower_limit)
```

```
Q1: 13.825000000000001
Q3: 17.175
IQR: 3.3499999999999996
upper limit: 22.2
lower limit: 8.8
```

```
In [14]: data['acceleration'] = np.where(data['acceleration']>upper_limit, upper_limit,
np.where(data['acceleration']<lower_limit,lower_limit,
data['acceleration']))
```

```
In [15]: #label encoding for name
data['origin'].value_counts()
```

```
Out[15]: 1    249
3     79
2     70
Name: origin, dtype: int64
```

```
In [16]: # categorical_features = ['origin','cylinders','model year' ]
```

```
In [17]: from sklearn.preprocessing import LabelEncoder
origin_enco = LabelEncoder()
data['origin'] = origin_enco.fit_transform(data['origin'])
```

```
In [18]: cylinder_enco = LabelEncoder()  
data['cylinders'] = cylinder_enco.fit_transform(data['cylinders'])
```

```
In [19]: data['cylinders'].value_counts()
```

```
Out[19]:  
1      204  
4      103  
3       84  
0        4  
2         3  
Name: cylinders, dtype: int64
```

```
In [20]: model_enco = LabelEncoder()  
data['model_year'] = model_enco.fit_transform(data['model_year'])
```

Feature Scaling

```
In [21]: data.drop("name",axis = 1, inplace = True)
```

```
In [22]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 398 entries, 0 to 397  
Data columns (total 8 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   mpg                   398 non-null   float64  
1   cylinders              398 non-null   int64  
2   displacement           398 non-null   float64  
3   weight                 398 non-null   int64  
4   acceleration           398 non-null   float64  
5   model_year            398 non-null   int64  
6   origin                 398 non-null   int64  
7   horsepower_new         398 non-null   int64  
dtypes: float64(3), int64(5)  
memory usage: 25.0 KB
```

```
In [23]: x = data.iloc[:,2:]  
y = data.iloc[:,0:1]
```

```
In [24]: x.head()
```

```
Out[24]:
```

	displacement	weight	acceleration	model_year	origin	horsepower_new
0	307.0	3504	12.0	0	0	130
1	350.0	3693	11.5	0	0	165
2	318.0	3436	11.0	0	0	150
3	304.0	3433	12.0	0	0	150
4	302.0	3449	10.5	0	0	140

```
In [25]: y.head()
```

```
Out[25]:
```

	mpg
0	18.0
1	15.0

2 18.0

3 16.0

4 17.0

```
In [26]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
sc_x = scaler.fit_transform(x)
pd.DataFrame(sc_x)
```

```
Out[26]:
```

	0	1	2	3	4	5
0	1.090604	0.630870	-1.320595	-1.627426	-0.715145	0.673589
1	1.503514	0.854333	-1.506489	-1.627426	-0.715145	1.590266
2	1.196232	0.550470	-1.692383	-1.627426	-0.715145	1.197404
3	1.061796	0.546923	-1.320595	-1.627426	-0.715145	1.197404
4	1.042591	0.565841	-1.878278	-1.627426	-0.715145	0.935497
...
393	-0.513026	-0.213324	0.017842	1.621983	-0.715145	-0.478804
394	-0.925936	-0.993671	2.471644	1.621983	0.533222	-1.369289
395	-0.561039	-0.798585	-1.469311	1.621983	-0.715145	-0.531185
396	-0.705077	-0.408411	1.133207	1.621983	-0.715145	-0.662139
397	-0.714680	-0.296088	1.430637	1.621983	-0.715145	-0.583567

398 rows × 6 columns

```
In [27]: pd.DataFrame(sc_x).describe()
```

```
Out[27]:
```

	0	1	2	3	4	5
count	3.980000e+02	3.980000e+02	3.980000e+02	3.980000e+02	3.980000e+02	3.980000e+02
mean	-1.785283e-17	-1.606755e-16	-1.071170e-16	2.142340e-16	-5.355850e-17	1.428227e-16
std	1.001259e+00	1.001259e+00	1.001259e+00	1.001259e+00	1.001259e+00	1.001259e+00
min	-1.204411e+00	-1.604943e+00	-2.510317e+00	-1.627426e+00	-7.151448e-01	-1.526434e+00
25%	-8.563178e-01	-8.828266e-01	-6.420819e-01	-8.150739e-01	-7.151448e-01	-7.407114e-01
50%	-4.314040e-01	-1.973624e-01	-1.933672e-02	-2.721449e-03	-7.151448e-01	-3.216593e-01
75%	6.584879e-01	7.538337e-01	6.034085e-01	8.096310e-01	5.332220e-01	5.426356e-01
max	2.511784e+00	2.565185e+00	2.471644e+00	1.621983e+00	1.781589e+00	3.292665e+00

```
In [28]: var = sc_x
var.shape
```

```
Out[28]: (398, 6)
```

Checking multicollinearity

```
In [29]: from statsmodels.stats.outliers_influence import variance_inflation_factor
var = sc_x
vif = pd.DataFrame()
vif['variance_inflation_factor'] = [variance_inflation_factor(var,i) for i in range(var.
vif['features'] = x.columns
```

```
In [30]: vif
```

```
Out[30]:
```

	variance_inflation_factor	features
0	12.178992	displacement
1	10.498106	weight
2	2.596991	acceleration
3	1.244622	model_year
4	1.729181	origin
5	9.432523	horsepower_new

Vif of displacement column is higher so we will remove it

```
In [31]: data.drop('displacement',axis = 1, inplace = True)
data.head()
```

```
Out[31]:
```

	mpg	cylinders	weight	acceleration	model_year	origin	horsepower_new
0	18.0	4	3504	12.0	0	0	130
1	15.0	4	3693	11.5	0	0	165
2	18.0	4	3436	11.0	0	0	150
3	16.0	4	3433	12.0	0	0	150
4	17.0	4	3449	10.5	0	0	140

```
In [32]: x1 = data.iloc[:,2:]
y1 = data.iloc[:,0:1]
```

```
In [33]: y1.head()
```

```
Out[33]:
```

	mpg
0	18.0
1	15.0
2	18.0
3	16.0
4	17.0

```
In [34]: scaler1 = StandardScaler()
sc_x1 = scaler1.fit_transform(x1)
pd.DataFrame(sc_x1)
```

```
Out[34]:
```

	0	1	2	3	4
--	---	---	---	---	---

0	0.630870	-1.320595	-1.627426	-0.715145	0.673589
1	0.854333	-1.506489	-1.627426	-0.715145	1.590266
2	0.550470	-1.692383	-1.627426	-0.715145	1.197404
3	0.546923	-1.320595	-1.627426	-0.715145	1.197404
4	0.565841	-1.878278	-1.627426	-0.715145	0.935497
...
393	-0.213324	0.017842	1.621983	-0.715145	-0.478804
394	-0.993671	2.471644	1.621983	0.533222	-1.369289
395	-0.798585	-1.469311	1.621983	-0.715145	-0.531185
396	-0.408411	1.133207	1.621983	-0.715145	-0.662139
397	-0.296088	1.430637	1.621983	-0.715145	-0.583567

398 rows × 5 columns

```
In [35]: var1 = sc_x1
var1.shape
```

```
Out[35]: (398, 5)
```

```
In [36]: vif1 = pd.DataFrame()
vif1['variance_inflation_factor1'] = [variance_inflation_factor(var1,i) for i in range(v
vif1['features'] = x1.columns
```


```
In [37]: vif1
```

```
Out[37]:
```

	variance_inflation_factor1	features
0	6.120806	weight
1	2.499726	acceleration
2	1.228353	model_year
3	1.538155	origin
4	8.699791	horsepower_new

Finding correlation

```
In [38]: plt.figure(figsize=(20,15))
sns.heatmap(data.corr(), annot = True, cmap='coolwarm')
plt.show()
```



```
In [39]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size =0.2,random_state=101)
```

Rergrssion model 1

```
In [40]: from sklearn.linear_model import LinearRegression
lm = LinearRegression()
lm.fit(x_train,y_train)
```

```
Out[40]: ▼ LinearRegression
LinearRegression()
```

```
In [41]: print(lm.coef_)
print('*****'*5)
print(lm.intercept_)

[[ 0.01520499 -0.00688301  0.0592133   0.78616418  1.16995869 -0.01721717]]
*****
[36.60074904]
```

```
In [42]: x_test.head()
```

```
Out[42]: displacement weight acceleration model_year origin horsepower_new
```

130	122.0	2451	16.5	4	0	80
202	258.0	3193	17.8	6	0	95
322	86.0	2110	17.9	10	2	65
104	400.0	4906	12.5	3	0	167
91	400.0	4464	12.0	3	0	150

Predict test dataset with linear model

```
In [43]: y_pred = lm.predict(x_test)
         y_pred
```

```
Out[43]: array([[24.32978998],
                [22.68151893],
                [33.52757722],
                [ 9.13806293],
                [12.44344079],
                [26.08186269],
                [34.43798179],
                [25.11370197],
                [27.03714153],
                [24.0735922 ],
                [25.75123137],
                [26.39348539],
                [34.75037624],
                [28.60873554],
                [17.20164285],
                [18.55360875],
                [20.71415173],
                [19.86982306],
                [25.66218644],
                [25.39752065],
                [ 8.58575882],
                [24.33951096],
                [29.39886739],
                [20.72518866],
                [15.50028393],
                [32.80409456],
                [25.35716315],
                [29.64592711],
                [17.42654862],
                [ 9.77994572],
                [20.60806978],
                [34.06511164],
                [24.67313184],
                [26.07496705],
                [25.8439045 ],
                [11.60777446],
                [28.3439356 ],
                [30.20284646],
                [15.95486995],
                [24.47684353],
                [32.71758896],
                [16.40550344],
                [26.66595053],
                [14.16428926],
                [21.65874182],
                [19.54208705],
                [29.05947031],
                [22.66456658],
```

```
[21.20134455],
[33.00662901],
[21.428808 ],
[22.16992788],
[23.22771654],
[25.69373339],
[15.66447169],
[27.83069754],
[34.78846079],
[25.32294323],
[16.80688139],
[32.01631988],
[30.40880856],
[25.43215002],
[35.08598643],
[19.09575745],
[21.40004846],
[23.79206428],
[31.05429952],
[32.55619221],
[12.68297108],
[13.22560415],
[18.92585729],
[25.0982099 ],
[19.39392958],
[27.9750004 ],
[18.9963028 ],
[12.11979723],
[25.45758 ],
[13.62522468],
[21.31933959],
[20.75414116]])
```

Evaluation for model 1

```
In [44]: from sklearn.metrics import r2_score
print("Accuracy: ", r2_score(y_test,y_pred))
```

```
Accuracy:  0.8050952184103181
```

linear model 2

```
In [45]: (x1_train,x1_test,y1_train,y1_test) = train_test_split(x1,y1,test_size = 0.2,random_stat
```

```
In [46]: lm1 = LinearRegression()
lm1.fit(x1_train,y1_train)
```

```
Out[46]: ▼ LinearRegression
LinearRegression()
```

```
In [47]: print(lm1.coef_)

[[-0.00565688 -0.01653479  0.77483059  0.93703182 -0.01014362]]
```

```
In [48]: print(lm1.intercept_)

[36.52142764]
```

```
In [49]: y1_pred = lm1.predict(x1_test)
```


y1_pred

```
Out[49]: array([[24.67141909],
 [21.85002448],
 [33.25246854],
 [ 9.19258794],
 [11.87363863],
 [26.84693733],
 [33.81272775],
 [25.41645485],
 [26.2725117 ],
 [24.34405765],
 [25.41768553],
 [26.84256453],
 [34.26322503],
 [28.16717412],
 [17.01226883],
 [18.17578991],
 [19.90359117],
 [19.45393895],
 [25.7879024 ],
 [25.80089887],
 [ 8.52586436],
 [24.30678357],
 [29.42283771],
 [20.20832814],
 [15.73158117],
 [32.06505479],
 [25.70529563],
 [29.84726814],
 [17.31077758],
 [ 9.0420631 ],
 [20.81716152],
 [33.79204695],
 [25.19650056],
 [26.51599687],
 [27.093569  ],
 [11.61432794],
 [28.5560865 ],
 [30.48753813],
 [15.75164895],
 [24.2109772 ],
 [32.8024156 ],
 [16.36215631],
 [26.55833956],
 [14.32627697],
 [21.58586971],
 [18.99067271],
 [28.9632257 ],
 [23.43148299],
 [20.89292508],
 [32.51698249],
 [20.59101696],
 [21.96319569],
 [22.63250839],
 [26.18607954],
 [15.41712019],
 [28.40804476],
 [34.40706377],
 [25.18803771],
 [16.9366266 ],
 [31.66258724],
 [30.35631524],
 [25.4684354 ],
 [34.73539456],
 [19.16741528],
```

```
[21.34044718],
[23.36415765],
[30.79613519],
[32.59075319],
[13.18962885],
[12.94689774],
[18.40691334],
[25.21394   ],
[19.27554839],
[28.3054246 ],
[16.61662767],
[12.6632652 ],
[25.78774242],
[13.5767135 ],
[20.76031012],
[20.64604912]])
```

```
In [50]: from sklearn.metrics import r2_score
```

```
In [51]: print("Accuracy: ", r2_score(y1_test, y1_pred))
```

```
Accuracy:  0.8078071902824695
```

OLS method

```
In [52]: # model1
```

```
from statsmodels.regression.linear_model import OLS
import statsmodels.regression.linear_model as smf
```

```
In [53]: reg1 = smf.OLS(endog = y_train, exog=x_train).fit()
```

```
In [54]: reg_test1 = smf.OLS(endog = y_test, exog = x_test).fit()
```

```
In [55]: reg1.summary()
```

```
Out[55]:
```

OLS Regression Results

Dep. Variable:	mpg	R-squared (uncentered):	0.967
Model:	OLS	Adj. R-squared (uncentered):	0.967
Method:	Least Squares	F-statistic:	1543.
Date:	Fri, 19 Jan 2024	Prob (F-statistic):	1.37e-228
Time:	22:25:17	Log-Likelihood:	-929.90
No. Observations:	318	AIC:	1872.
Df Residuals:	312	BIC:	1894.
Df Model:	6		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
displacement	0.0137	0.009	1.573	0.117	-0.003	0.031
weight	-0.0088	0.001	-8.521	0.000	-0.011	-0.007
acceleration	1.6214	0.084	19.239	0.000	1.456	1.787
model_year	1.0404	0.074	14.139	0.000	0.896	1.185

origin	2.1056	0.404	5.210	0.000	1.310	2.901
horsepower_new	0.1359	0.016	8.257	0.000	0.103	0.168
Omnibus:	1.485	Durbin-Watson:	2.004			
Prob(Omnibus):	0.476	Jarque-Bera (JB):	1.533			
Skew:	0.163	Prob(JB):	0.465			
Kurtosis:	2.900	Cond. No.	4.91e+03			

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[3] The condition number is large, 4.91e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [56]: `reg_test1.summary()`

Out[56]:

OLS Regression Results						
Dep. Variable:	mpg	R-squared (uncentered):	0.977			
Model:	OLS	Adj. R-squared (uncentered):	0.975			
Method:	Least Squares		F-statistic:	514.5		
Date:	Fri, 19 Jan 2024		Prob (F-statistic):	3.30e-58		
Time:	22:25:17		Log-Likelihood:	-217.81		
No. Observations:	80		AIC:	447.6		
Df Residuals:	74		BIC:	461.9		
Df Model:	6					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
displacement	-0.0257	0.015	-1.760	0.083	-0.055	0.003
weight	-0.0041	0.001	-2.962	0.004	-0.007	-0.001
acceleration	1.4029	0.117	11.946	0.000	1.169	1.637
model_year	0.6845	0.135	5.073	0.000	0.416	0.953
origin	3.0225	0.748	4.043	0.000	1.533	4.512
horsepower_new	0.1234	0.028	4.431	0.000	0.068	0.179
Omnibus:	3.085	Durbin-Watson:	2.136			
Prob(Omnibus):	0.214	Jarque-Bera (JB):	2.425			
Skew:	0.406	Prob(JB):	0.297			
Kurtosis:	3.258	Cond. No.	5.44e+03			

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[3] The condition number is large, 5.44e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [57]: reg2 = smf.OLS(endog=y1_train,exog = x1_train).fit()
```

```
In [58]: reg2_test = smf.OLS(endog = y1_test,exog=x1_test).fit()
```

```
In [59]: reg2.summary()
```

Out[59]:

OLS Regression Results						
Dep. Variable:	mpg		R-squared (uncentered):	0.967		
Model:	OLS		Adj. R-squared (uncentered):	0.967		
Method:	Least Squares		F-statistic:	1842.		
Date:	Fri, 19 Jan 2024		Prob (F-statistic):	1.04e-229		
Time:	22:25:17		Log-Likelihood:	-931.15		
No. Observations:	318		AIC:	1872.		
Df Residuals:	313		BIC:	1891.		
Df Model:	5					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
weight	-0.0077	0.001	-10.132	0.000	-0.009	-0.006
acceleration	1.5502	0.071	21.747	0.000	1.410	1.691
model_year	1.0298	0.073	14.021	0.000	0.885	1.174
origin	1.8944	0.382	4.958	0.000	1.143	2.646
horsepower_new	0.1419	0.016	8.851	0.000	0.110	0.173
Omnibus:	0.757	Durbin-Watson:	1.984			
Prob(Omnibus):	0.685	Jarque-Bera (JB):	0.750			
Skew:	0.117	Prob(JB):	0.687			
Kurtosis:	2.959	Cond. No.	4.64e+03			

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[3] The condition number is large, 4.64e+03. This might indicate that there are strong multicollinearity or other numerical problems.


```
In [60]: plt.scatter(y_test,y_pred)
plt.show()
```




```
In [61]: plt.scatter(y1_test,y1_pred)
plt.show()
```



```
In [62]: sns.distplot((y_test-y_pred),bins = 50)
plt.show()
```



```
In [63]: sns.distplot((y1_test-y1_pred),bins = 50)
plt.show()
```



```
In [64]: import statsmodels.api as sm
from statsmodels.formula.api import ols
```

```
In [65]: data1 = pd.read_csv('mpg.csv')
data1.head()
```

```
Out[65]:   mpg  cylinders  displacement  horsepower  weight  acceleration  model_year  origin  name
```

0	18.0	8	307.0	130	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	ford torino

```
In [66]: model = ols('mpg~displacement', data=data1).fit()
annova_result = sm.stats.anova_lm(model, typ=2)
print(annova_result)
```

	sum_sq	df	F	PR(>F)
displacement	15685.163618	1.0	724.994303	1.655889e-91
Residual	8567.411859	396.0	NaN	NaN

```
In [67]: data.head()
```

Out[67]:

	mpg	cylinders	weight	acceleration	model_year	origin	horsepower_new
0	18.0	4	3504	12.0	0	0	130
1	15.0	4	3693	11.5	0	0	165
2	18.0	4	3436	11.0	0	0	150
3	16.0	4	3433	12.0	0	0	150
4	17.0	4	3449	10.5	0	0	140